

Alpha-Beta Cutoff


Outline

I. Alpha-beta pruning


II. Alpha-beta pruning algorithm

I. Alpha-Beta Cutoff

- ♣ #states is exponential in the depth of the game tree.
- ♦ But we can often compute the correct minimax decision by **pruning large parts of the tree** that do not affect the outcome.



Node Pruning



The player will not move to node n if it has a **better choice**

- ♦ either at the same level (e.g., m')
- ♦ or at any node (e.g., m) higher up in the tree.

Prune n once we have found enough about it to reach the above conclusion.

Alpha and Beta Values

Alpha-beta pruning gets its name from two extra parameters α, β

α = the *highest*-value (i.e., the best choice) so far along a path for the player MAX. // MAX's best score so far.

Alpha and Beta Values

Alpha-beta pruning gets its name from two extra parameters α, β

α = the *highest*-value (i.e., the best choice) so far along a path for the player MAX. // MAX's best score so far.

eventual value $\geq \alpha$ “at least”

Alpha and Beta Values

Alpha-beta pruning gets its name from two extra parameters α, β

α = the *highest*-value (i.e., the best choice) so far along a path for the player MAX. // MAX's best score so far.

eventual value $\geq \alpha$ “at least”

β = the *lowest*-value (i.e., the best choice) so far along a path for the player MIN. // MIN's best score so far.

eventual value $\leq \beta$ “at most”

Alpha and Beta Values

Alpha-beta pruning gets its name from two extra parameters α, β

α = the *highest*-value (i.e., the best choice) so far along a path for the player MAX. // MAX's best score so far.

eventual value $\geq \alpha$ “at least”


β = the *lowest*-value (i.e., the best choice) so far along a path for the player MIN. // MIN's best score so far.

eventual value $\leq \beta$ “at most”

- ◆ Update the values of α and β as the search goes along.
- ◆ Prune the remaining branches
 - at a MIN node with current value $\leq \alpha$ (because its parent, a MAX node, has from another child or is passed on the value α already).
 - at a MAX node with current value $\geq \beta$ (because its parent, a MIN node, has or is passed on the value β already).


Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.





Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.





Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.





Re-examining the Game Tree

Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Re-examining the Game Tree


Fig. 5.5 in the 4th edition of the textbook *incorrectly executes* the algorithm ALPHA-BETA-SEARCH in Fig. 5.7.




Cont'd




Cont'd




Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd




Cont'd




Cont'd




Cont'd




Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd





Cont'd




Cont'd




Cont'd




$$\begin{aligned}
 \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\
 &= \max(3, \min(2, x, y), 2) \\
 &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\
 &= 3.
 \end{aligned}$$




A Larger Example (Wikipedia)

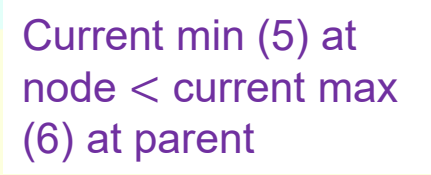


A Larger Example (Wikipedia)




A Larger Example (Wikipedia)






A Larger Example (Wikipedia)



A Larger Example (Wikipedia)



Exercise (suggested): Trace the execution of ALPHA-BETA-SEARCH (next page) and dynamically update the $[\alpha, \beta]$ value of every expanded node.

II. Alpha-Beta Search Algorithm (3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm (3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm (3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm (3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v // maximum value of all children if none are pruned



function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v

II. Alpha-Beta Search Algorithm (3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

// no change of β value within MAX-VALUE()

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v *// pruning (returned v will not affect parent's β value).*

$\alpha \leftarrow \text{MAX}(\alpha, v)$ *// $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.*

return v *// maximum value of all children if none are pruned*



function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$

for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v

$\beta \leftarrow \text{MIN}(\beta, v)$

return v

II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

// no change of β value within MAX-VALUE()

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v // maximum value of all children if none are pruned



function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v // pruning (no change will occur to parent's α value).

$\beta \leftarrow \text{MIN}(\beta, v)$

return v



II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

// no change of β value within MAX-VALUE()

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v // maximum value of all children if none are pruned



function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v // pruning (no change will occur to parent's α value).

$\beta \leftarrow \text{MIN}(\beta, v)$ // $v > \alpha$: new β passed on to the rest MIN-VALUE calls within the for loop.

return v



II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

// no change of β value within MAX-VALUE()

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v // maximum value of all children if none are pruned



function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v // pruning (no change will occur to parent's α value).

$\beta \leftarrow \text{MIN}(\beta, v)$ // $v > \alpha$: new β passed on to the rest MIN-VALUE calls within the for loop.

return v // minimum value of all children if none are pruned.



II. Alpha-Beta Search Algorithm

(3rd Edition of Textbook)

function ALPHA-BETA-SEARCH(*state*) **returns** an action

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in $\text{ACTIONS}(\text{state})$ with value v

// no change of β value within MAX-VALUE()

function MAX-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow -\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \geq \beta$ **then return** v // pruning (returned v will not affect parent's β value).

$\alpha \leftarrow \text{MAX}(\alpha, v)$ // $v < \beta$: new α passed on to the remaining MIN-VALUE calls within the for loop.

return v // maximum value of all children if none are pruned



// no change of α value within MIN-VALUE()

function MIN-VALUE (*state*, α , β) **returns** a utility value

if $\text{TERMINAL-TEST}(\text{state})$ **then return** $\text{UTILITY}(\text{state})$

$v \leftarrow +\infty$


for each a in $\text{ACTIONS}(\text{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$

if $v \leq \alpha$ **then return** v // pruning (no change will occur to parent's α value).

$\beta \leftarrow \text{MIN}(\beta, v)$ // $v > \alpha$: new β passed on to the rest MIN-VALUE calls within the for loop.

return v // minimum value of all children if none are pruned.



Short Summary on the Algorithm

Both MAX and MIN nodes

- ◆ Receive α and β values from the parent.
- ◆ Pass the best value of the executed actions back to the parent.



MAX node

- ◆ Updates the α value if one of its actions yields a better value.
- ◆ Does not update the β value.
- ◆ Use the β value to skip actions that do not affect the outcome.

MIN node



- ◆ Updates the β value if one of its actions yields a better value.
- ◆ Does not update the α value.
- ◆ Use the α value to skip actions that do not affect the outcome.

Move Ordering



Successors 14 and 5 would've been pruned had 2 been generated first.



Move Ordering



Successors 14 and 5 would've been pruned had 2 been generated first.

- Effectiveness of pruning is highly dependent on the order in which successors are generated.



Move Ordering



Successors 14 and 5 would've been pruned had 2 been generated first.

- Effectiveness of pruning is highly dependent on the order in which successors are generated.
- “Perfect ordering” has effective branching factor $\sqrt{b}$, which limits examination to only $O(b^{m/2})$ nodes compared to $O(b^m)$ for minimax.

Move Ordering



Successors 14 and 5 would've been pruned had 2 been generated first.

- Effectiveness of pruning is highly dependent on the order in which successors are generated.
- “Perfect ordering” has effective branching factor $\sqrt{b}$, which limits examination to only $O(b^{m/2})$ nodes compared to $O(b^m)$ for minimax.
- $O(b^{3m/4})$ nodes for random move ordering.

Two Strategies


- ♦ For chess, a simple ordering function (sequentially considering captures, threats, forward moves, backward moves) could get close to the best case.
- ♠ Even with alpha-beta pruning and clever move ordering, minimax won't work well enough for games like chess and Go due to their vast state spaces.

Two Strategies

- ♦ For chess, a simple ordering function (sequentially considering captures, threats, forward moves, backward moves) could get close to the best case.
- ♠ Even with alpha-beta pruning and clever move ordering, minimax won't work well enough for games like chess and Go due to their vast state spaces.

Claude Shannon (1950) suggested two strategies:

- Type A ($\Rightarrow$ heuristic alpha-beta tree search) – chess
 - ♣ Considers all possible moves to a certain depth.
 - ♣ Use a heuristic function to estimate utilities of states at that depth.




Explores wide & shallow portion of the search tree.

Two Strategies

- ◆ For chess, a simple ordering function (sequentially considering captures, threats, forward moves, backward moves) could get close to the best case.
- ♠ Even with alpha-beta pruning and clever move ordering, minimax won't work well enough for games like chess and Go due to their vast state spaces.

Claude Shannon (1950) suggested two strategies:

- Type A ($\Rightarrow$ heuristic alpha-beta tree search) – chess
 - ♣ Considers all possible moves to a certain depth.
 - ♣ Use a heuristic function to estimate utilities of states at that depth.
- Type B ($\Rightarrow$ Monte Carlo tree search) – Go
 - ♣ Ignore moves that look bad.
 - ♣ Follow promising lines “as far as possible”.



Explores wide & shallow portion of the search tree.



Explores deep but narrow portion of the search tree.